

Using monolingual
data

Review: the noisy channel

Bayes' rule

$$p(\mathbf{e} \mid \mathbf{f}) = \frac{p_{LM}(\mathbf{e}) \times p_{TM}(\mathbf{f} \mid \mathbf{e})}{p(\mathbf{f})}$$

Review: the noisy channel

chain rule

$$p(\mathbf{e}, \mathbf{f}) = p_{LM}(\mathbf{e}) \times p_{TM}(\mathbf{f} \mid \mathbf{e})$$

Review: the noisy channel

chain rule

$$p(\mathbf{e}, \mathbf{f}) = p_{LM}(\mathbf{e}) \times p_{TM}(\mathbf{f} | \mathbf{e})$$

language model

translation model
(conditional
language model)

How much training data for each of these?

Review: the noisy channel

chain rule


$$p(\mathbf{e}, \mathbf{f}) = p_{LM}(\mathbf{e}) \times p_{TM}(\mathbf{f} | \mathbf{e})$$

language model

translation model
(conditional
language model)

maybe 1B words

Impact on size of language model training data (in words) on quality of Arabic-English statistical machine translation system




Review: neural MT

Er, there's no decomposition here.

$$p(\mathbf{e} \mid \mathbf{f}) = p(\mathbf{e} \mid \mathbf{f})$$

language model

translation model
(conditional
language model)



Where did
this go?

Is this data
still useful?

maybe 1B words

Attempt 1

Gulcehre et al., June 2015

- MaxEnt (NN mumbo-jumbo term: “shallow fusion”)

$$\log p(\mathbf{e} \mid \mathbf{f}) = \log p_{TM}(\mathbf{e} \mid \mathbf{f}) + \beta \log p_{LM}(\mathbf{e})$$

Attempt 1

Gulcehre et al., June 2015

- MaxEnt (NN mumbo-jumbo: “shallow fusion”)

$$\log p(\mathbf{e} \mid \mathbf{f}) = \log p_{TM}(\mathbf{e} \mid \mathbf{f}) + \beta \log p_{LM}(\mathbf{e})$$



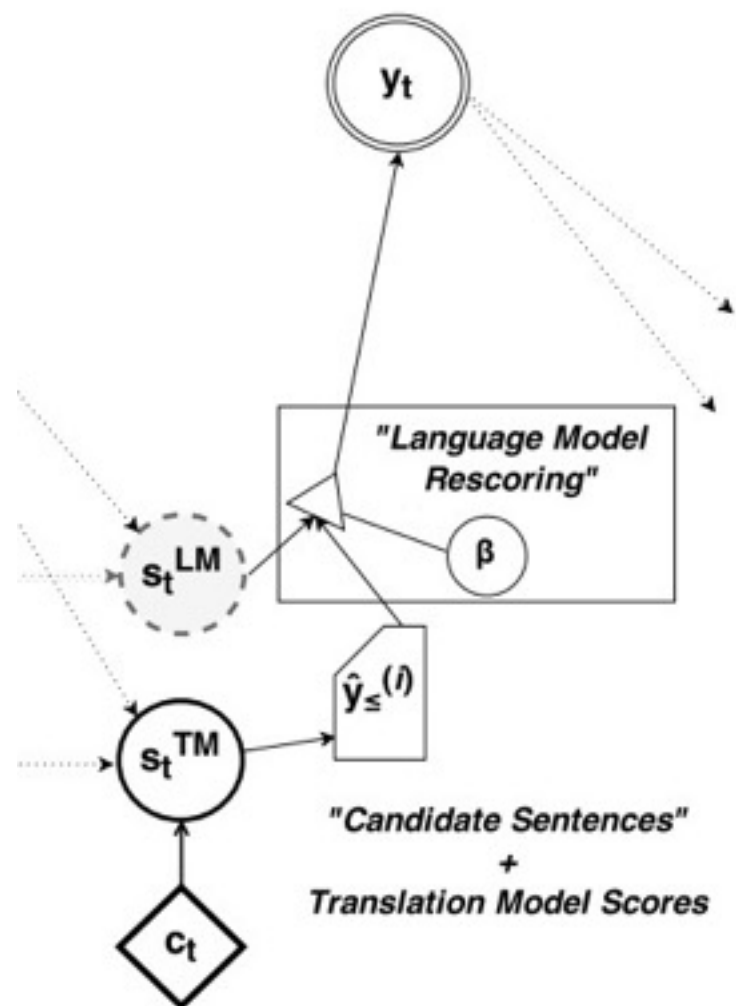
Tune this parameter on
development set, once
other models are learned

Note difference from Bayes' rule!


Attempt 1

Gulcehre et al., June 2015

- “Deep fusion”: output distribution a function of (fixed) LM hidden state and (learned) TM hidden state



(a) Shallow Fusion (Sec. 4.1)



(b) Deep Fusion (Sec. 4.2)

Attempt 1

Gulcehre et al., June 2015

- “Deep fusion”: output distribution a function of (fixed) LM hidden state and (learned) TM hidden state

	De-En		Cs-En	
	Dev	Test	Dev	Test
NMT Baseline	25.51	23.61	21.47	21.89
Shallow Fusion	25.53	23.69	21.95	22.18
Deep Fusion	25.88	24.00	22.49	22.36

	Development Set		Test Set			
	dev2010	tst2010	tst2011	tst2012	tst2013	Test 2014
Previous Best (Single)	15.33	17.14	18.77	18.62	18.88	-
Previous Best (Combination)	-	17.34	18.83	18.93	18.70	-
NMT	14.50	18.01	18.40	18.77	19.86	18.64
NMT+LM (Shallow)	14.44	17.99	18.48	18.80	19.87	18.66
NMT+LM (Deep)	15.69	19.34	20.17	20.23	21.34	20.56

Attempt 2

Sennrich et al., Nov 2015

- Idea 1: Train encoder-decoder on parallel and monolingual target-language data.
- On monolingual data, set context to 0 (essentially, train decoder without encoder).

Attempt 2

Sennrich et al., Nov 2015

- Idea 2: Very old idea: backtranslation
 - MT folklore: in English->Russian->English:
“The spirit is willing, but the flesh is weak” ->
“The vodka is good, but the meat is rotten”
- Basic idea: train two systems, backtranslate English to French, then retrain English French-English system on resulting data.

Attempt 2

Sennrich et al., Nov 2015

- Monolingual: Empty context vector
- Synthetic: Backtranslation

name	training instances	BLEU			
		newstest2014		newstest2015	
		single	ens-4	single	ens-4
syntax-based (Sennrich and Haddow, 2015)		22.6	-	24.4	-
Neural MT (Jean et al., 2015b)		-	-	22.4	-
parallel	37m (parallel)	19.9	20.4	22.8	23.6
+monolingual	49m (parallel) / 49m (monolingual)	20.4	21.4	23.2	24.6
+synthetic	44m (parallel) / 36m (synthetic)	22.7	23.8	25.7	26.5

Table 3: English→German translation performance (BLEU) on WMT training/test sets. Ens-4: ensemble of 4 models. Number of training instances varies due to differences in training time and speed.

Attempt 2

Sennrich et al., Nov 2015

- Synthetic: Backtranslation

name	BLEU	
	2014	2015
PBSMT (Haddow et al., 2015)	28.8	29.3
NMT (Gülçehre et al., 2015)	23.6	-
+shallow fusion	23.7	-
+deep fusion	24.0	-
parallel	25.9	26.7
+synthetic	29.5	30.4
+synthetic (ensemble of 4)	30.8	31.6

Table 5: German→English translation performance (BLEU) on WMT training/test sets (newstest2014; newstest2015).

Attempt 2

Sennrich et al., Nov 2015

- Synthetic: Backtranslation

system	BLEU	
	WMT	IWSLT
parallel	20.1	21.5
+synthetic	20.8	21.6
PBSMT gain	+0.7	+0.1
NMT gain	+2.9	+1.2

phrase-based MT {

Table 8: Phrase-based SMT results (English→German) on WMT test sets (average of newstest201{4,5}), and IWSLT test sets (average of tst201{3,4,5}), and average BLEU gain from adding synthetic data for both PBSMT and NMT.

Attempt 2

Sennrich et al., Nov 2015

- Synthetic: Backtranslation

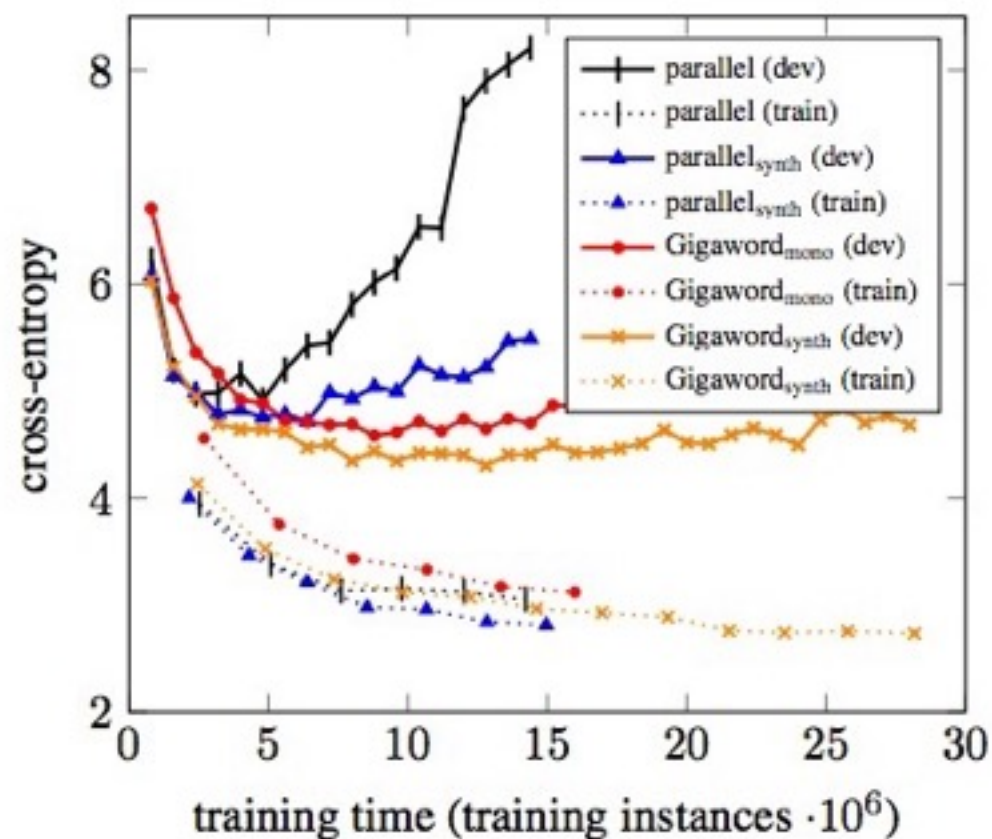


Figure 1: Turkish→English training and development set (tst2010) cross-entropy as a function of training time (number of training instances) for different systems.

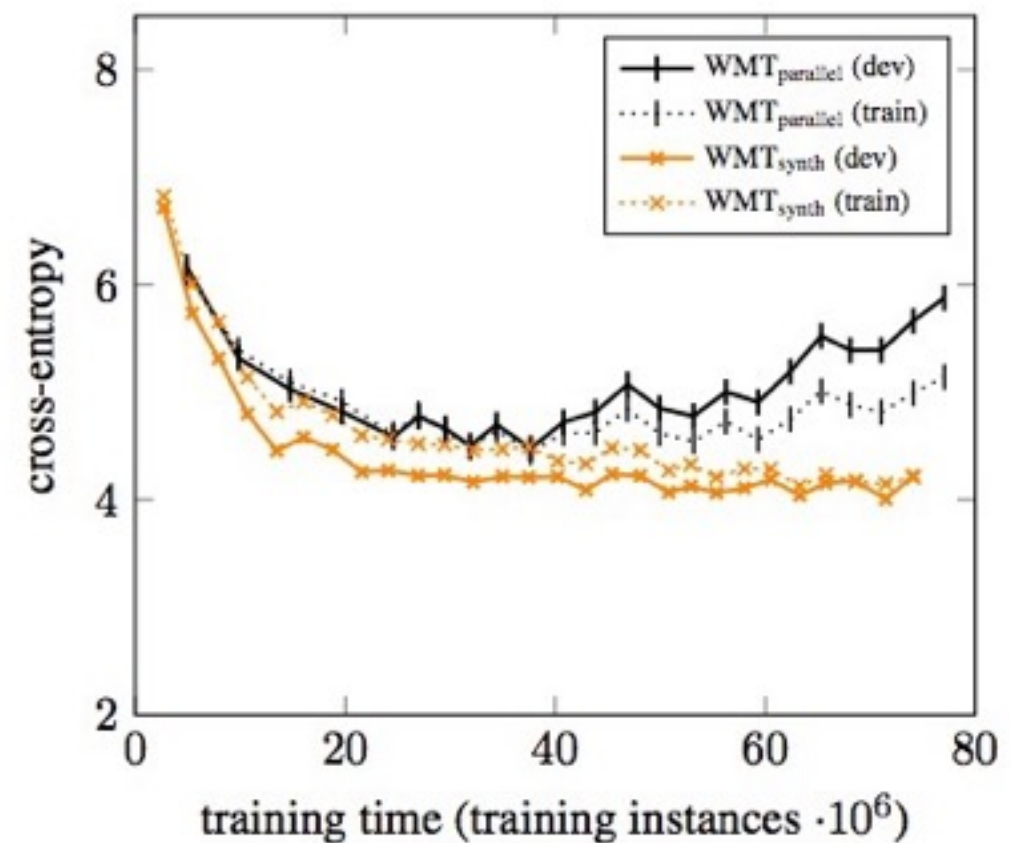


Figure 2: English→German training and development set (newstest2013) cross-entropy as a function of training time (number of training instances) for different systems.

Summary: LMs in NMT

- Monday's discussion:
 - Phrase-based MT may be better for **adequacy**.
 - Neural MT may be better for **fluency**.
- Do we need a LM if our MT is already fluent?
- More generally: if our discriminative model has enough capacity, do we need a strong prior?
- On the other hand: denoising the input seems to work!

Hybrid systems

- Monday's discussion:
 - Phrase-based MT may be better for **adequacy**.
 - Neural MT may be better for **fluency**.
- How can we combine these benefits?

Simpler: LMs

Q: How would we combine a neural
and an n-gram language model?

Simpler: LMs

Q: How would we combine a neural and an n-gram language model?

interpolation of k LMs:

$$P(w_i|\mathbf{c}) = \sum_{k=1}^K \lambda_k(\mathbf{c}) P_k(w_i|\mathbf{c})$$

k th weight
(weights must sum to 1)

k th LM

Simpler: LMs

Neubig & Dyer 2016

$$P(w_i|\mathbf{c}) = \sum_{k=1}^K \lambda_k(\mathbf{c}) P_k(w_i|\mathbf{c})$$

interpolation of k LMs as matrix multiplication:

Probabilities \mathbf{p}^\top

Coefficients $\boldsymbol{\lambda}^\top$

$$\begin{bmatrix} p_1 \\ p_2 \\ \vdots \\ p_J \end{bmatrix} = \underbrace{\begin{bmatrix} d_{1,1} & d_{1,2} & \cdots & d_{1,K} \\ d_{2,1} & d_{2,2} & \cdots & d_{2,K} \\ \vdots & \vdots & \ddots & \vdots \\ d_{J,1} & d_{J,2} & \cdots & d_{J,K} \end{bmatrix}}_{\text{Distribution matrix } D} \begin{bmatrix} \lambda_1 \\ \lambda_2 \\ \vdots \\ \lambda_K \end{bmatrix}$$

k th weight
(weights must sum to 1)

k th LM

Simpler: LMs

Neubig & Dyer 2016

$$P(w_i|\mathbf{c}) = \sum_{k=1}^K \lambda_k(\mathbf{c}) P_k(w_i|\mathbf{c})$$

interpolation of k LMs as matrix multiplication:

Probabilities \mathbf{p}^\top

Coefficients $\boldsymbol{\lambda}^\top$

$$\begin{bmatrix} p_1 \\ p_2 \\ \vdots \\ p_J \end{bmatrix} = \begin{bmatrix} d_{1,1} & d_{1,2} & \cdots & d_{1,K} \\ d_{2,1} & d_{2,2} & \cdots & d_{2,K} \\ \vdots & \vdots & \ddots & \vdots \\ d_{J,1} & d_{J,2} & \cdots & d_{J,K} \end{bmatrix} \begin{bmatrix} \lambda_1 \\ \lambda_2 \\ \vdots \\ \lambda_K \end{bmatrix}$$

$\underbrace{\hspace{10em}}$
Distribution matrix D

probability of
 i th element

Simpler: LMs

Neubig & Dyer 2016

$$P(w_i|\mathbf{c}) = \sum_{k=1}^K \lambda_k(\mathbf{c}) P_k(w_i|\mathbf{c})$$

heuristic **interpolation** of LMs (e.g. Kneser-Ney)

$$\begin{array}{cc} \text{Probabilities } \mathbf{p}^\top & \text{Heuristic interp. coefficients } \boldsymbol{\lambda}^\top \\ \underbrace{\begin{bmatrix} p_1 \\ p_2 \\ \vdots \\ p_J \end{bmatrix}} & = \underbrace{\begin{bmatrix} d_{1,1} & d_{1,2} & \cdots & d_{1,N} \\ d_{2,1} & d_{2,2} & \cdots & d_{2,N} \\ \vdots & \vdots & \ddots & \vdots \\ d_{J,1} & d_{J,2} & \cdots & d_{J,N} \end{bmatrix}}_{\text{Count-based probabilities } P_C(w_i = j | w_{i-n+1}^{i-1})} \underbrace{\begin{bmatrix} \lambda_1 \\ \lambda_2 \\ \vdots \\ \lambda_N \end{bmatrix}} \end{array}$$

Simpler: LMs

Neubig & Dyer 2016

Probabilities \mathbf{p}^\top

Heuristic interp. coefficients $\boldsymbol{\lambda}^\top$

$$\begin{bmatrix} p_1 \\ p_2 \\ \vdots \\ p_J \end{bmatrix} = \begin{bmatrix} d_{1,1} & d_{1,2} & \cdots & d_{1,N} \\ d_{2,1} & d_{2,2} & \cdots & d_{2,N} \\ \vdots & \vdots & \ddots & \vdots \\ d_{J,1} & d_{J,2} & \cdots & d_{J,N} \end{bmatrix} \begin{bmatrix} \lambda_1 \\ \lambda_2 \\ \vdots \\ \lambda_N \end{bmatrix}$$

Do we know any other way to produce this vector?

Count-based probabilities $P_C(w_i = j | w_{i-n+1}^{i-1})$

Simpler: LMs

Neubig & Dyer 2016

Probabilities p^\top

Result of $\text{softmax}(\text{NN}(\mathbf{c}))$

$$\begin{bmatrix} p_1 \\ p_2 \\ \vdots \\ p_J \end{bmatrix} = \begin{bmatrix} d_{1,1} & d_{1,2} & \cdots & d_{1,N} \\ d_{2,1} & d_{2,2} & \cdots & d_{2,N} \\ \vdots & \vdots & \ddots & \vdots \\ d_{J,1} & d_{J,2} & \cdots & d_{J,N} \end{bmatrix} \begin{bmatrix} \lambda_1 \\ \lambda_2 \\ \vdots \\ \lambda_N \end{bmatrix}$$

Count-based probabilities $P_C(w_i = j | w_{i-n+1}^{i-1})$

Simpler: LMs

Neubig & Dyer 2016

$$P(w_i|\mathbf{c}) = \sum_{k=1}^K \lambda_k(\mathbf{c}) P_k(w_i|\mathbf{c})$$

matrix **interpretation** of RNNLM

Probabilities \mathbf{p}^\top Result of softmax(NN(\mathbf{c}))

$$\begin{bmatrix} p_1 \\ p_2 \\ \vdots \\ p_J \end{bmatrix} = \begin{bmatrix} 1 & 0 & \cdots & 0 \\ 0 & 1 & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & 1 \end{bmatrix} \begin{bmatrix} \lambda_1 \\ \lambda_2 \\ \vdots \\ \lambda_J \end{bmatrix}$$

$\underbrace{\hspace{10em}}$
J-by-J identity matrix I

Simpler: LMs

Neubig & Dyer 2016

$$P(w_i|\mathbf{c}) = \sum_{k=1}^K \lambda_k(\mathbf{c}) P_k(w_i|\mathbf{c})$$

matrix **interpretation** of RNNLM

Probabilities \mathbf{p}^\top Result of softmax(NN(\mathbf{c}))

$$\begin{bmatrix} p_1 \\ p_2 \\ \vdots \\ p_J \end{bmatrix} = \begin{bmatrix} 1 & 0 & \cdots & 0 \\ 0 & 1 & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & 1 \end{bmatrix} \begin{bmatrix} \lambda_1 \\ \lambda_2 \\ \vdots \\ \lambda_J \end{bmatrix}$$

$\underbrace{\hspace{10em}}$
J-by-J identity matrix I

Simpler: LMs

Neubig & Dyer 2016

$$P(w_i|\mathbf{c}) = \sum_{k=1}^K \lambda_k(\mathbf{c}) P_k(w_i|\mathbf{c})$$

neural interpolation of count-based and RNNLM

Probabilities \mathbf{p}^\top

Result of softmax(NN(\mathbf{c}))

$$\begin{bmatrix} p_1 \\ p_2 \\ \vdots \\ p_J \end{bmatrix} = \underbrace{\begin{bmatrix} d_{1,1} & \cdots & d_{1,N} & 1 & \cdots & 0 \\ d_{2,1} & \cdots & d_{2,N} & 0 & \cdots & 0 \\ \vdots & \ddots & \vdots & \vdots & \ddots & \vdots \\ d_{J,1} & \cdots & d_{J,N} & 0 & \cdots & 1 \end{bmatrix}}_{\text{Count-based probabilities and } J\text{-by-}J \text{ identity matrix}} \begin{bmatrix} \lambda_1 \\ \lambda_2 \\ \vdots \\ \lambda_{J+N} \end{bmatrix}$$

Count-based probabilities and J -by- J identity matrix

Simpler: LMs

Neubig & Dyer 2016

$$P(w_i|\mathbf{c}) = \sum_{k=1}^K \lambda_k(\mathbf{c}) P_k(w_i|\mathbf{c})$$

matrix **interpretation** of RNNLM

Probabilities \mathbf{p}^\top Result of softmax(NN(\mathbf{c}))

$$\begin{bmatrix} p_1 \\ p_2 \\ \vdots \\ p_J \end{bmatrix} = \begin{bmatrix} 1 & 0 & \cdots & 0 \\ 0 & 1 & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & 1 \end{bmatrix} \begin{bmatrix} \lambda_1 \\ \lambda_2 \\ \vdots \\ \lambda_J \end{bmatrix}$$

$\underbrace{\hspace{10em}}$
J-by-J identity matrix I

Simpler: LMs

Neubig & Dyer 2016

- Training: block dropout (force count-based interpolation weights to zero for first several epochs)

	Dist.	Interp.	PPL
(1)	KN	HEUR	140.8/156.5
(2)	δ	LSTM	105.9/116.9
(3)	KN	LSTM	135.2/149.1
(4)	KN, δ	LSTM -BIDO	108.4/130.4
(5)	KN, δ	LSTM +BIDO	95.3 /104.5

Simpler: LMs

Neubig & Dyer 2016

- Training: block dropout (force count-based interpolation weights to zero for first several epochs).
- Other observations: seem to help with low-frequency words (where count-based LMs tend to be better)

Applied to MT

Arthur et al. 2016

- Observation: phrase-based MT better at adequacy (often due to handling of rare words).

Input:	I come from <u>Tunisia</u> .
Reference:	<u>チュニジア</u> の出身です。 <u>Chunisia</u> no shusshindesu. (<i>I'm from Tunisia.</i>)
System:	<u>ノルウェー</u> の出身です。 <u>Noruue-</u> no shusshindesu. (<i>I'm from Norway.</i>)

Figure 1: An example of a mistake made by NMT on low-frequency content words.

Applied to MT

Arthur et al. 2016

- Observation: phrase-based MT better at adequacy (often due to handling of rare words).

Input:	I come from <u>Tunisia</u> .
Reference:	<u>チュニジア</u> の出身です。 <u>Chunisia</u> no shusshindesu. (<i>I'm from Tunisia.</i>)
System:	<u>ノルウェー</u> の出身です。 <u>Noruue-</u> no shusshindesu. (<i>I'm from Norway.</i>)

Figure 1: An example of a mistake made by NMT on low-frequency content words.

Applied to MT

Arthur et al. 2016

- Observation: phrase-based MT better at adequacy (often due to handling of rare words).
- Basic idea: interpolate conditional neural LM with IBM Model 1 probabilities.

$$p_o(e_i|F, e_1^{i-1}) = \begin{bmatrix} p_l(e_i = 1|F, e_1^{i-1}) & p_m(e = 1|F, e_1^{i-1}) \\ \vdots & \vdots \\ p_l(e_i = |V_e||F, e_1^{i-1}) & p_m(e = |V_e||F, e_1^{i-1}) \end{bmatrix} \begin{bmatrix} \lambda \\ 1 - \lambda \end{bmatrix}$$

Applied to MT

Arthur et al. 2016

- Observation: phrase-based MT better at adequacy (often due to handling of rare words).
- Basic idea: interpolate conditional neural LM with IBM Model 1 probabilities.

System	<u>BTEC</u>			<u>KFTT</u>		
	BLEU	NIST	RECALL	BLEU	NIST	RECALL
pbmt	48.18	6.05	27.03	22.62	5.79	13.88
hiero	52.27	6.34	24.32	22.54	5.82	12.83
attn	48.31	5.98	17.39	20.86	5.15	17.68
auto-bias	49.74*	6.11*	50.00	23.20[†]	5.59[†]	19.32
hyb-bias	50.34[†]	6.10*	41.67	22.80[†]	5.55[†]	16.67

Where to next?

- Coursework 2 due Monday (right before lecture).
- You've seen many things that are likely to be useful to you in Coursework 3.
- Next two weeks: more advanced ideas
Including modeling more language phenomena
- Week 10: final lectures (won't need to understand these for coursework).